On Linear Logic, Functional Programming, and Attack Trees

Harley Eades III[1], Jiaming Jiang[2], and Aubrey Bryant[1]

[1] Computer Science, Augusta University, harley.eades@gmail.com
[2] Computer Science, North Carolina State University

Abstract. TODO

1 Introduction

Attack trees are perhaps the most popular graphical model used to conduct threat analysis of both physical and virtual secure systems. They were made popular by Bruce Schneier in the late nineties [15]. In those early years attack trees were studied and used as a syntactic tool to help guide analysis. However, as systems grew more complex the need for a semantics of attack trees become apparent, after all, without a proper semantics how can we safely manipulate attack trees, extend the expressivity of attack trees, or compare them?

A number of different models of attack trees have been proposed: a model in boolean algebras [11,10,14], series-parallel pomsets [12], Petri nets [13], and tree automata [1]. There have also been various extensions, such as, adding sequential composition [6], and defense nodes [9,10]. All of these models and extensions have their benefits, but at the heart of them all is logic.

The model in boolean algebras was the first and most elegant model of attack trees, but it failed to capture the process notion of attack trees, that is, the fact that base attacks are actual processes that need to be carried out, and the branching nodes composed these processes in different ways. Thus, the community moved towards models of resources like parallel-series pomsets, Petri nets, and automata. However, the complexity of these models increased, and hence, comparing the models is difficult which makes it hard to decide which to use and under which circumstances, and so we have wondered, is there a means of recovering the elegant model in logic, and can this logic teach us anything new?

Linear Logic. It is fitting that attack trees are the most popular model used in threat analysis, because one of the most widely studied logics used to reason about resources is *linear logic* which is an excellent candidate for modeling attack trees. In fact, Horne et al.[5] has already produced a number of interesting results. Most importantly, they show that attack trees can be modeled as formulas in linear logic, and then one can prove properties between attack trees by proving implications between them. Furthermore, by studying attack trees from a linear logical perspective they introduce a new property between attack trees called *specializations*. Prior to their paper the literature was primarily concerned with equality between attack trees, but implication allows us to break that equality

up. A specialization is a one way relationship between two attack trees that may even be influenced by the attack trees attribute domain; we will discuss specializations in more detail in Section **??**.

This paper has two main contributions, the first is a new simple linear logical semantics of causal attack trees – attack trees with sequential composition – in four-valued truth tables. We show that our semantics is surprisingly expressive. It supports specializations, and even lays outside of the semantics proposed by Horne et al.[5], because it combines in an interesting way what they call the *ideal* and *filter* semantics of causal attack trees. However, the most appealing aspect of this semantics is that it is extremely simple. Furthermore, we introduce a family of natural deduction systems based in the logic of bunched implications which were formed by studying the truth table semantics.

Functional Programming. Our second contribution is Lina, a new domain specific functional programming language for conducting threat analysis using attack trees. Consider the example attack trees in Fig. 1. Both of these are

A.

```
seq_node "ATM attack"
 (and_node "get credientials"
   (base_na "steal card")
   (or_node "get PIN"
    (base_na "social engineer")
    (base_na "find a post-it")))
 (base_na "withdraw money")
```

B.

```
or_node "ATM attack"
 (seq_node "attack vector 1"
    (and_node "get credentials 1"
       (base_na "social engineer")
       (base_na "steal card"))
    (base_na "withdraw money"))
 (seq_node "attack vector 2"
    (and_node "get credentials 2"
       (base_na "steal card")
       (base_na "find a post-it"))
    (base_na "withdraw money"))
```

Fig. 1. Attack Tree for an ATM attack from Figure 1 (A) and Figure 2 (B) of Kordy et al. [8]

actual Lina programs, in fact every example in this paper are Lina programs. Lina supports causal attack trees both with attributes or without, thus, there are two types of base attacks: base attacks with attributes, denoted `base_wa`, and base attacks with no attributes, denoted `base_na`. Lina is designed to be extremely simple, and actually reflect the typical pseudocode found throughout the literature. However, Lina is more than just a simple definitional language.

Lina is an embedded domain specific language whose host language is the Haskell programming language [7]. So, why purely functional and why strongly typed? As security researchers/professionals we are in the business of verifying the correctness of various systems. Thus, our tools should be taking advantage of verification tools to insure that our constructions, tools, and analysis are correct. By embedding Lina into Haskell we are able to take advantage of cutting edge verification tools while conducting threat analysis. For example, right out the box

Lina supports using property-based randomized testing using QuickCheck [2], and refinement types in Liquid Haskell [16] to verify properties of our attack trees or the attribute domains used while analyzing attack trees. Furthermore, strong typing helps catch bugs while we develop our attack trees and their attribute domains as a side-effect of type checking. Finally, functional programs are short, but not obfuscated, and hence, allow for very compact and trustworthy programs.

That being said, we are designing Lina so that it can be used with very little Haskell experience. It is our hope that one will be able to make use of Lina without having to know how to write Haskell programs, and we plan to develop new tooling to support this.

Lina is approaching threat analysis from a programming language perspective. This approach leads to a number of new advances. First, as Gadyatskaya and Trujillo-Rasua [4] argue as a community we need to start building more automated means of conducting threat analysis, and there is no better way to build or connect automated tools than a programming language. Lina is perfect as a target for new tools, and it can be connected to existing tools fairly easily. In fact, Lina already supports automation using the automatic rewrite system Maude [3], for example, the two attack trees in Fig. 1 can be automatically proven equivalent to each other in Lina. This is similar to Krody's [8] SPTool, but Lina goes further and supports more than one backend rewrite system, for example, Lina is the first tool to support automatically proving specializations of attack trees. The user can choose which backend they wish to use.

We have a number of extensions planned for the future like supporting attack-defense trees, attack(-defense) graphs, and attack nets. We plan to support even more automation using SAT and SMT solvers. It is our hope that Lina grows into a one stop shop for threat analysis.

2 Causal Attack Trees

We begin by introducing causal attack trees. This formulation of attack trees was first proposed by Jhawar et al. [6] where they called them causal attack trees, but sequential composition does not always meet the same properties as conjunction, for example, classically it is a self dual operator, thus, we follow Horne et al.'s lead [5] and call them causal attack trees.

Definition 1. *Suppose* B *is a set of base attacks whose elements are denoted by* b*. Then an* ***attack tree*** *is defined by the following grammar:*

$$A, B, C, T := b \mid \mathsf{OR}(A, B) \mid \mathsf{AND}(A, B) \mid \mathsf{SEQ}(A, B)$$

Equivalence of attack trees, denoted by $A \approx B$*, is defined as follows:*

$$\mathsf{OR}(A, A) \approx A$$
$$\mathsf{OR}(\mathsf{OR}(A, B), C) \approx \mathsf{OR}(A, \mathsf{OR}(B, C))$$
$$\mathsf{AND}(\mathsf{AND}(A, B), C) \approx \mathsf{AND}(A, \mathsf{AND}(B, C))$$
$$\mathsf{SEQ}(\mathsf{SEQ}(A, B), C) \approx \mathsf{SEQ}(A, \mathsf{SEQ}(B, C))$$
$$\mathsf{OR}(A, B) \approx \mathsf{OR}(B, A)$$
$$\mathsf{AND}(A, B) \approx \mathsf{AND}(B, A)$$
$$\mathsf{AND}(A, \mathsf{OR}(B, C)) \approx \mathsf{OR}(\mathsf{AND}(A, B), \mathsf{AND}(A, C))$$
$$\mathsf{SEQ}(A, \mathsf{OR}(B, C)) \approx \mathsf{OR}(\mathsf{SEQ}(A, B), \mathsf{SEQ}(A, C))$$

This definition of causal attack trees differs slightly from Jhawar et. al.'s [6] definition. They define n-ary operators, but we only consider the binary case, because it fits better with the models presented here and it does not loose any generality because we can model the n-ary case using binary operators in the obvious way.

Throughout the sequel we will show that the previous rules are sound with respect to our new models, but just as Horne et al. [5] we will then show that there are properties of attack trees that these rules do not support.

3 A Quaternary Semantics for Causal Attack Trees

Kordy et al. [10] gave a very elegant and simple semantics of attack-defense trees in boolean algebras. Unfortunately, while their semantics is elegant it does not capture the resource aspect of attack trees, it allows contraction, and it does not provide a means to model sequential composition. In this section we give a semantics of attack trees in the spirit of Kordy et al.'s using a four valued logic.

The propositional variables of our quaternary logic, denoted by A, B, C, and D, range over the set $\mathsf{4} = \{0, \frac{1}{4}, \frac{1}{2}, 1\}$. We think of 0 and 1 as we usually do in boolean algebras, but we think of $\frac{1}{4}$ and $\frac{1}{2}$ as intermediate values that can be used to break various structural rules. In particular we will use these values to prevent exchange for sequential composition from holding, and contraction from holding for parallel and sequential composition.

Definition 2. *The logical connectives of our four valued logic are defined as follows:*

Parallel and Sequential Conjunction:

$$A \odot_4 B = 1, \text{ where neither } A \text{ nor } B \text{ are } 0$$
$$A \odot_4 B = 0, \text{otherwise}$$

$$A \rhd_4 B = 1, \text{ where } A \in \{\tfrac{1}{2}, 1\} \text{ and } B \neq 0$$
$$\tfrac{1}{4} \rhd_4 B = \tfrac{1}{4}, \text{where } B \neq 0$$
$$A \rhd_4 B = 0, \text{otherwise}$$

Choice: $A \sqcup_4 B = \mathsf{max}(A, B)$

These definitions are carefully crafted to satisfy the necessary properties to model attack trees. Comparing these definitions with Kordy et al.'s [10] work we can see that choice is defined similarly, but parallel composition is not a product –

ordinary conjunction – but rather a linear tensor product, and sequential composition is not actually definable in a boolean algebra, and hence, makes heavy use of the intermediate values to insure that neither exchange nor contraction hold.

We use the usual notion of equivalence between propositions, that is, propositions ϕ and ψ are considered equivalent, denoted by $\phi \equiv \psi$, if and only if they have the same truth tables. In order to model attack trees the previously defined logical connectives must satisfy the appropriate equivalences corresponding to the equations between attack trees. These equivalences are all proven by the following result.

Lemma 1 (Properties of the Attack Tree Operators in the Quaternary Semantics).

(Symmetry) For any A and B, $A \bullet B \equiv B \bullet A$, for $\bullet \in \{\odot_4, \sqcup_4\}$.

(Symmetry for Sequential Conjunction) It is not the case that, for any A and B, $A \rhd_4 B \equiv B \rhd_4 A$.

(Associativity) For any A, B, and C, $(A \bullet B) \bullet C \equiv A \bullet (B \bullet C)$, for $\bullet \in \{\odot_4, \rhd_4, \sqcup_4\}$.

(Contraction for Parallel and Sequential Conjunction) It is not the case that for any A, $A \bullet A \equiv A$, for $\bullet \in \{\odot_4, \rhd_4\}$.

(Distributive Law) For any A, B, and C, $A \bullet (B \sqcup_4 C) \equiv (A \bullet B) \sqcup_4 (A \bullet C)$, for $\bullet \in \{\odot_4, \rhd_4\}$.

Proof. Symmetry, associativity, contraction for choice, and the distributive law for each operator hold by simply comparing truth tables. As for contraction for parallel composition, suppose $A = \frac{1}{4}$. Then by definition $A \odot_4 A = 1$, but $\frac{1}{4}$ is not 1. Contraction for sequential composition also fails, suppose $A = \frac{1}{2}$. Then by definition $A \rhd_4 A = 1$, but $\frac{1}{2}$ is not 1. Similarly, symmetry fails for sequential composition. Suppose $A = \frac{1}{4}$ and $B = \frac{1}{2}$. Then $A \rhd_4 B = \frac{1}{4}$, but $B \rhd_4 A = 1$.

At this point it is quite easy to model attack trees as formulas. The following defines their interpretation.

Definition 3. *Suppose $\mathbb{B}$ is some set of base attacks, and $\nu : \mathbb{B} \to \mathsf{PVar}$ is an assignment of base attacks to propositional variables. Then we define the interpretation of attack trees to propositions as follows:*

$$\begin{array}{lcllcl} [\![b \in \mathbb{B}]\!] & = & \nu(b) & [\![\mathsf{SEQ}(A,B)]\!] & = & [\![A]\!] \rhd_4 [\![B]\!] \\ [\![\mathsf{AND}(A,B)]\!] & = & [\![A]\!] \odot_4 [\![B]\!] & [\![\mathsf{OR}(A,B)]\!] & = & [\![A]\!] \sqcup_4 [\![B]\!] \end{array}$$

We can use this semantics to prove equivalences between attack trees.

Lemma 2 (Equivalence of Attack Trees in the Quaternary Semantics). *Suppose $\mathbb{B}$ is some set of base attacks, and $\nu : \mathbb{B} \to \mathsf{PVar}$ is an assignment of base attacks to propositional variables. Then for any attack trees A and B, if $A \approx B$, then $[\![A]\!] \equiv [\![B]\!]$.*

Proof. This proof holds by induction on the form of $A \approx B$.

This is a very simple and elegant semantics, but it also leads to a more substantial theory.

4 Specialization in the Quaternary Semantics

The quaternary semantics introduced in the previous section does indeed capture all of the equivalences of attack trees, but it also supports proving specializations. Consider the example attack trees in Fig. 2. Attack tree C is a sound

A.

```
and_node "obtain secret"
 (or_node "obtain encryped file"
    (base_na "bribe sysadmin")
    (base_na "steal backup"))
 (seq_node "obtain password"
    (base_na "break into system")
    (base_na "install keylogger"))
```

B.

```
seq_node "break in, obtain secret"
 (base_na "break into system")
 (and_node "obtain secret inside"
   (base_na "install keylogger")
   (base_na "steal backup"))
```

C.

```
or_node "obtain secret"
 (and_node "obtain secret via sysadmin"
  (base_na "bribe sysadmin")
  (seq_node "obtain password"
    (base_na "break into system")
    (base_na "install keylogger")))
  (seq_node "break in, obtain secrect"
    (base_na "break into system")
    (and_node "obtain secret inside"
      (base_na "install keylogger")
      (base_na "steal backup")))
```

Fig. 2. Encrypted Data Attack from Figure 1 (A), Figure 3 (B), and Figure 2 (C) of Horne et al. [5]

specialization of attack tree A, and attack tree B is a sound specialization of attack tree A. Attack tree C requires the attacker to break into the system before they can steal the backup, but attack tree A does not require this. Then attack tree B has dropped bribing the sysadmin and simply requires the attacker to just steal the backups. Notice that none of the attack trees in Fig. 2 are equivalent. So how do we prove these specializations are sound?

We simply define a notion of entailment in the quaternary semantics. Denote by $A \leqslant_4 B$ the obvious ordering on $\mathbf{4}$. Then we have the following result immediately.

Lemma 3 (Entailment in the Quaternary Semantics). *$A \equiv B$ if and only if $A \leqslant_4 B$ and $B \leqslant_4 A$*

This result shows that we can now break up the equivalence of attack trees into directional properties captured here by entailments, and hence, every equivalence proved in the previous section can also be used directionally.

We may now formally define when a attack tree is a sound specialization of another attack tree.

Definition 4. *An attack tree A is a sound specialization of an attack B if and only if $[\![A]\!] \leqslant_4 [\![B]\!]$.*

The next result proves some additional properties in the quaternary semantics that can be used to reason about attack trees.

Lemma 4 (Properties of Entailment in the Quaternary Semantics).

Ideal Entailments:

$((a \odot_4 b) \rhd_4 (c \odot_4 d)) \leqslant_4 ((a \rhd_4 c) \odot_4 (b \rhd_4 d))$
$((a \odot_4 b) \rhd_4 c) \leqslant_4 (a \odot_4 (b \rhd_4 c))$
$(a \rhd_4 (b \odot_4 c)) \leqslant_4 (b \odot_4 (a \rhd_4 c))$
$(a \rhd_4 b) \leqslant_4 (a \odot_4 b)$

Filter Entailments:

$((a \rhd_4 c) \odot_4 (b \rhd_4 d)) \leqslant_4 ((a \odot_4 b) \rhd_4 (c \odot_4 d))$
$(a \odot_4 (b \rhd_4 c)) \leqslant_4 ((a \odot_4 b) \rhd_4 c)$

Choice Entailments:

$a \leqslant_4 (a \sqcup_4 b)$
$b \leqslant_4 (a \sqcup_4 b)$

Each of the above entailments are due to Horne et al. [5]. They introduce two types of semantics called the *ideal semantics* and the *filter semantics*. The former satisfies all of the entailments in Lemma 1 and the left side of Lemma 4, but the latter is similar, however, satisfying the right side of Lemma 4. They were able to show that the ideal entailments allow one to prove properties of attack trees with different attribute domains than the filter entailments.

In comparison with their work the semantics presented here is a blend of the ideal and filter semantics. It primarily consists of the ideal semantics, but as we can see from Lemma 4 the first two axioms are actually logical equivalences. This implies that this semantics can prove properties that neither the ideal nor the filter semantics can capture. However, we do not yet know which attribute domains correspond to this semantics. We can isolate ourselves to just the ideal or the filter semantics, but what attribute domains can we reason about in this semantics when we blend them? This is left to future work.

We can now formally prove that the attack tree C is a specialization of attack trees A and B in Fig. 2.

Example 1. First, consider the following assignment:

$a :=$ `"bribe sysadmin"` $\quad b :=$ `"break into system"`
$c :=$ `"install keylogger"` $\quad d :=$ `"steal backup"`

Then we have the following interpretations:

$$
\begin{aligned}
[\![A]\!] &= [\![\mathsf{AND}(\mathsf{OR}(a,d),\mathsf{SEQ}(b,c))]\!] \\
&= (a \sqcup_4 d) \odot_4 (b \rhd_4 c)
\end{aligned}
\qquad
\begin{aligned}
[\![B]\!] &= [\![\mathsf{SEQ}(b,\mathsf{AND}(c,d))]\!] \\
&= b \rhd_4 (c \odot_4 d)
\end{aligned}
$$

$$
\begin{aligned}
[\![C]\!] &= [\![\mathsf{OR}(\mathsf{AND}(a,\mathsf{SEQ}(b,c)),\mathsf{SEQ}(b,\mathsf{AND}(c,d)))]\!] \\
&= (a \odot_4 (b \rhd_4 c)) \sqcup_4 (b \rhd_4 (c \odot_4 d))
\end{aligned}
$$

We reuse the same names for base attacks across the interpretations above. Finally, we have the following two entailments:

$$
\begin{aligned}
&\underline{[\![C]\!] \leqslant_4 [\![A]\!]}: \\
&\quad (a \odot_4 (b \rhd_4 c)) \sqcup_4 (b \rhd_4 (c \odot_4 d)) \\
&\leqslant_4 (a \odot_4 (b \rhd_4 c)) \sqcup_4 (b \rhd_4 (d \odot_4 c)) \\
&\leqslant_4 (a \odot_4 (b \rhd_4 c)) \sqcup_4 (d \odot_4 (b \rhd_4 c)) \\
&\leqslant_4 (a \sqcup_4 d) \odot_4 (b \rhd_4 c)
\end{aligned}
\qquad
\begin{aligned}
&\underline{[\![B]\!] \leqslant_4 [\![A]\!]}: \\
&\quad b \rhd_4 (c \odot_4 d) \\
&\leqslant_4 b \rhd_4 (c \odot_4 (a \sqcup_4 d)) \\
&\leqslant_4 b \rhd_4 ((a \sqcup_4 d) \odot_4 c) \\
&\leqslant_4 (a \sqcup_4 d) \odot_4 (b \rhd_4 c)
\end{aligned}
$$

Notice that neither $A \leqslant_4 C$ nor $A \leqslant_4 B$ hold, and thus, equivalences cannot prove the previous properties.

5 Lina: An EDSL for Conducting Threat Analysis using Causal Attack Trees

All of the models mentioned in this paper have been incorporated into a new embedded domain specific language (EDSL) for conducting threat analysis called Lina; which means small, young palm tree, but we constructed the name by combining the words linear and attack.

Lina is embedded inside of Haskell, a statically typed purely functional programming language. The most important property of any EDSL is that they subsume the entirety of their host language, and can be prototyped quite rapidly. Haskell has several advantages, like Lina's ability to utilize Haskell's cutting edge verification tools, and its strong type system for catching bugs quickly. In addition, Haskell has several tools that make building EDSLs more easily, for example, type classes.

Lina currently supports three types of causal attack trees:

- Process Attack Trees: these are attack trees with no attributes at all,
- Attributed Process Attack Trees: these are attack trees with attributes on the base attacks only. This is an intermediate representation used to build full attack trees.
- Full Attack Trees: these are attributed process attack trees with an associated attribute domain.

Internally, we represent causal attack trees by a simple data type:

```
data IAT where
  Base :: ID -> IAT
  OR   :: ID -> IAT -> IAT -> IAT
  AND  :: ID -> IAT -> IAT -> IAT
  SEQ  :: ID -> IAT -> IAT -> IAT
```

where ID is a type synonym of Integer. We then define each type of attack tree:

```
-- Attributed Process Attack Tree
data APAttackTree attribute label =
 APAttackTree {
  process_tree :: IAT,
  labels :: B.Bimap label ID,
  attributes :: M.Map ID attribute
 }
```

```
-- Process Attack Tree
type PAttackTree label = APAttackTree () label

-- Full Attack Tree
data AttackTree attribute label = AttackTree {
      ap_tree :: APAttackTree attribute label,
      configuration :: Conf attribute
}
```

A B.Bimap is a dictionary where we can efficiently lookup IDs given a `label` or efficiently lookup `label`s given an ID, a M.Map is a typical dictionary, and () is the unit type.

The previous data types reveal that actually all attack trees are attributed process attack trees where a process attack tree simply does not use the attributes. This design has two main benefits: internal attack trees are very easy to translate to various backends, especially formulas because we can use the IDs on base attacks as atomic formulas – which has its own benefits discussed below – and the second benefit is that modifying labels and attributes is more efficient than having them labeled on the trees themselves.

One important aspect of the definition of the various forms of attack trees is that the types `label` and `attribute` are actually type variables, and thus, our definition of attack trees is very general, in fact, `label` and `attribute` can be instantiated with any type whose elements are comparable. This property is captured by ad-hoc polymorphism using type classes in Haskell, and are checked during type checking.

Conducting threat analysis using attack trees requires them to be associated with an attribute domain. Typically, an attribute domain is a set together with operations for computing the attribute of the branching nodes of an attack tree given attributes on the base attacks. In Lina attribute domains are defined by a type, here called `attribute`, and a configuration:

```
data Conf attribute = (Ord attribute) => Conf {
      andOp :: attribute -> attribute -> attribute,
      seqOp :: attribute -> attribute -> attribute
}
```

Utilizing higher-order functions we can define configurations quite easily, and quite generically. For example, here is the configuration that computes the max attribute for parallel nodes and takes the sum of the children of sequential nodes as the attribute for sequential nodes:

```
maxAddConf :: (Ord attribute,Semiring attribute) => Conf attribute
maxAddConf = Conf max (.+.)
```

Notice here that this configuration will work with any type at all whose elements are comparable and form a semiring. This includes types like Integer and Double.

The definitional language for attributed process attack trees of type `APAttackTree attribute label` is described by the following grammar:

$$
\begin{aligned}
at ::= & \texttt{base\_na label} \\
& \mid \texttt{base\_wa attribute label} \\
& \mid \texttt{or\_node label at1 at2} \\
& \mid \texttt{and\_node label at1 at2} \\
& \mid \texttt{seq\_node label at1 at2}
\end{aligned}
$$

A full example of the definition of an attributed process attack tree for attacking an autonomous vehicle can be found in Fig. 5. The definition of `vehicle_attack`

```
vehicle_attack :: APAttackTree Double String
vehicle_attack = start_PAT $
  or_node "Autonomous Vehicle Attack"
    (seq_node "external sensor attack"
       (base_wa 0.2 "modify street signs to cause wreck")
       (and_node "social engineering attack"
          (base_wa 0.6 "pose as mechanic")
          (base_wa 0.1 "install malware")))
    (seq_node "over night attack"
       (base_wa 0.05 "Find address where car is stored")
       (seq_node "compromise vehicle"
          (or_node "break in"
             (base_wa 0.8 "break window")
             (base_wa 0.5 "disable door alarm/locks"))
          (base_wa 0.1 "install malware")))
```

Fig. 3. Lina Script for an Autonomous Vehicle Attack

begins with a call to `start_PAT`. Behind the scenes all of the `ID`'s within the internal attack tree are managed implicitly, and this requires the internals of Lina to work within a special state-based type. The function `start_PAT` initializes this state.

Finally, we can define the vehicle attack tree as follows:

```
vehicle_AT :: AttackTree Double String
vehicle_AT = AttackTree vehicle_attack maxMaxConf
```

This attack tree associates the vehicle attack attributed process attack tree with a configuration called `maxMaxConf` that simply takes the maximum as the attribute of every parallel and sequential node.

Two features that Lina has that other tools lack is its ability to abstract the definitions of attack trees, and it is highly compositional, because it is embedded inside of a functional programming language. Consider the following abstraction of `vehicle_attack`:

```
vehicle_AT' :: Conf Double -> AttackTree Double String
vehicle_AT' conf = AttackTree vehicle_attack conf
```

Here the configuration has been abstracted. This facilitates experimentation because the security practitioner can run several different forms of analysis on the same attack tree using different attribute domains.

Attack trees in Lina can also be composed, and hence, complex trees can be broken down into smaller ones, then studied in isolation. This helps facilitate

correctness, and offers more flexibility. As an example, in Fig. 4 we break up `vehicle_attack` into several smaller attack trees. We can see in the example

```
se_attack :: APAttackTree Double String
se_attack = start_PAT $
  and_node "social engineering attack"
     (base_wa 0.6 "pose as mechanic")
     (base_wa 0.1 "install malware")
```

```
bi_attack :: APAttackTree Double String
bi_attack = start_PAT $
  or_node "break in"
     (base_wa 0.8 "break window")
     (base_wa 0.5 "disable door alarm/locks")
```

```
cv_attack :: APAttackTree Double String
cv_attack = start_PAT $
  seq_node "compromise vehicle"
    (insert bi_attack)
    (base_wa 0.1 "install malware")
```

```
es_attack :: APAttackTree Double String
es_attack = start_PAT $
  seq_node "external sensor attack"
       (base_wa 0.2 "modify street signs to cause
                     wreck")
       (insert se_attack)
```

```
on_attack :: APAttackTree Double String
on_attack = start_PAT $
  seq_node "over night attack"
     (base_wa 0.05 "Find address where car
                    is stored")
     (insert cv_attack)
```

```
vehicle_attack'' :: APAttackTree Double String
vehicle_attack'' = start_PAT $
  or_node "Autonomous Vehicle Attack"
    (insert es_attack)
    (insert on_attack)
```

Fig. 4. The Autonomous Vehicle Attack Decomposed

that if we wish to use an already defined attack tree in one we are defining, then we can make use of `insert`. Behind the scenes Lina maintains a special state that tracks the identifiers of each node, thus, when one wishes to insert an existing attack tree, which will have its own identifier labeling, into a new tree, then that internal state must be updated, thus, `insert` carries out this updating. Lina is designed so that the user never has to encounter that internal state.

References

1. S.A. Camtepe and B. Yener. Modeling and detection of complex attacks. In *Security and Privacy in Communications Networks and the Workshops, 2007. SecureComm 2007. Third International Conference on*, pages 234–243, Sept 2007.
2. Koen Claessen and John Hughes. Quickcheck: A lightweight tool for random testing of haskell programs. *SIGPLAN Not.*, 46(4):53–64, May 2011.
3. Manuel Clavel, Francisco Durán, Steven Eker, Patrick Lincoln, Narciso Martı-Oliet, José Meseguer, and Carolyn Talcott. Maude manual (version 2.1). *SRI International, Menlo Park*, 2005.
4. Olga Gadyatskaya and Rolando Trujillo-Rasua. New directions in attack tree research: Catching up with industrial needs. In Peng Liu, Sjouke Mauw, and Ketil Stolen, editors, *Graphical Models for Security*, pages 115–126, Cham, 2018. Springer International Publishing.
5. Ross Horne, Sjouke Mauw, and Alwen Tiu. Semantics for specialising attack trees based on linear logic. *Fundamenta Informaticae*, 153(1-2):57–86, 2017.

```
OR("Autonomous Vehicle Attack",0.6)
  (SEQ("external sensor attack",0.6)
     ("modify street signs to cause wreck",0.2)
     (AND("social engineering attack",0.6)
        ("pose as mechanic",0.6)
        ("install malware",0.1)))
```


```
OR("Autonomous Vehicle Attack",0.6)
  (SEQ("external sensor attack",0.6)
     ("modify street signs to cause wreck",0.2)
     (AND("social engineering attack",0.6)
        ("pose as mechanic",0.6)
        ("install malware",0.1)))
```



```
OR("Autonomous Vehicle Attack",0.6)
  (SEQ("external sensor attack",0.6)
     ("modify street signs to cause wreck",0.2)
     (AND("social engineering attack",0.6)
        ("pose as mechanic",0.6)
        ("install malware",0.1)))
```


Fig. 5. Set of Possible Attacks for an Autonomous Vehicle Attack

6. Ravi Jhawar, Barbara Kordy, Sjouke Mauw, Saša Radomirović, and Rolando Trujillo-Rasua. Attack trees with sequential conjunction. In Hannes Federrath and Dieter Gollmann, editors, *ICT Systems Security and Privacy Protection*, volume 455 of *IFIP Advances in Information and Communication Technology*, pages 339–353. Springer International Publishing, 2015.
7. Simon Peyton Jones. *Haskell 98 language and libraries: the revised report*. Cambridge University Press, 2003.
8. Barbara Kordy, Piotr Kordy, and Yoann van den Boom. *SPTool – Equivalence Checker for* `SAND` *Attack Trees*, pages 105–113. Springer International Publishing, Cham, 2017.
9. Barbara Kordy, Sjouke Mauw, Saša Radomirović, and Patrick Schweitzer. Foundations of attack–defense trees. In Pierpaolo Degano, Sandro Etalle, and Joshua Guttman, editors, *Formal Aspects of Security and Trust*, pages 80–95, Berlin, Heidelberg, 2011. Springer Berlin Heidelberg.
10. Barbara Kordy, Marc Pouly, and Patrick Schweitzer. Computational aspects of attack–defense trees. In Pascal Bouvry, MieczysławA. Kłopotek, Franck Leprévost, Małgorzata Marciniak, Agnieszka Mykowiecka, and Henryk Rybiński, editors, *Security and Intelligent Information Systems*, volume 7053 of *Lecture Notes in Computer Science*, pages 103–116. Springer Berlin Heidelberg, 2012.
11. Barbara Kordy, Marc Pouly, and Patrick Schweitzer. A probabilistic framework for security scenarios with dependent actions. In Elvira Albert and Emil Sekerinski, editors, *Integrated Formal Methods*, volume 8739 of *Lecture Notes in Computer Science*, pages 256–271. Springer International Publishing, 2014.
12. Sjouke Mauw and Martijn Oostdijk. Foundations of attack trees. In DongHo Won and Seungjoo Kim, editors, *Information Security and Cryptology - ICISC 2005*, volume 3935 of *Lecture Notes in Computer Science*, pages 186–198. Springer Berlin Heidelberg, 2006.

13. J. P. McDermott. Attack net penetration testing. In *Proceedings of the 2000 Workshop on New Security Paradigms*, NSPW '00, pages 15–21, New York, NY, USA, 2000. ACM.
14. L. Piètre-Cambacédès and M. Bouissou. Beyond attack trees: Dynamic security modeling with boolean logic driven markov processes (bdmp). In *Dependable Computing Conference (EDCC), 2010 European*, pages 199–208, April 2010.
15. Bruce Schneier. Attack trees: Modeling security threats. *Dr. Dobb's journal*, December 1999.
16. Niki Vazou, Eric L. Seidel, Ranjit Jhala, Dimitrios Vytiniotis, and Simon Peyton-Jones. Refinement types for haskell. *SIGPLAN Not.*, 49(9):269–282, August 2014.